

```

1 import numpy as np
2 import pandas as pd
3 import matplotlib.pyplot as plt
4
5 from sklearn.model_selection import train_test_split
6

--NORMAL--

1 from sklearn.metrics import mean_squared_error, r2_score
2 #Import Data from Input
3 df= pd.read_csv(r"C:\\Users\\shiny\\Downloads\\Automobile_data.csv")

```

```
1 df.head()
```

	symboling	normalized-losses	make	fuel-type	aspiration	num-of-doors	body-style	drive-wheels	engine-location
0	3	?	alfa-romero	gas	std	two	convertible	rwd	front
1	3	?	alfa-romero	gas	std	two	convertible	rwd	front
2	1	?	alfa-romero	gas	std	two	hatchback	rwd	front
3	2	164	audi	gas	std	four	sedan	fwd	front
4	2	164	audi	gas	std	four	sedan	4wd	front

```
1 df.shape
```

```
2
```

```
(205, 26)
```

```
1 df.columns
```

```

Index(['symboling', 'normalized-losses', 'make', 'fuel-type', 'aspiration',
       'num-of-doors', 'body-style', 'drive-wheels', 'engine-location',
       'wheel-base', 'length', 'width', 'height', 'curb-weight', 'engine-type',
       'num-of-cylinders', 'engine-size', 'fuel-system', 'bore', 'stroke',
       'compression-ratio', 'horsepower', 'peak-rpm', 'city-mpg',
       'highway-mpg', 'price'],
      dtype='object')

```

```
1 df.info()
```

```

<class 'pandas.core.frame.DataFrame'>
RangeIndex: 205 entries, 0 to 204
Data columns (total 26 columns):
 #   Column           Non-Null Count  Dtype  
--- 
 0   symboling        205 non-null    int64  
 1   normalized-losses 205 non-null    object  
 2   make             205 non-null    object  
 3   fuel-type        205 non-null    object  
 4   aspiration       205 non-null    object  
 5   num-of-doors     205 non-null    object  
 6   body-style       205 non-null    object  
 7   drive-wheels    205 non-null    object  
 8   engine-location  205 non-null    object  
 9   wheel-base       205 non-null    float64 
 10  length           205 non-null    float64 
 11  width            205 non-null    float64 
 12  height           205 non-null    float64 
 13  curb-weight      205 non-null    int64  
 14  engine-type      205 non-null    object  
 15  num-of-cylinders 205 non-null    object  
 16  engine-size      205 non-null    int64  
 17  fuel-system      205 non-null    object  
 18  bore             205 non-null    object  
 19  stroke           205 non-null    object  

```

```

20 compression-ratio 205 non-null float64
21 horsepower 205 non-null object
22 peak-rpm 205 non-null object
23 city-mpg 205 non-null int64
24 highway-mpg 205 non-null int64
25 price 205 non-null object
dtypes: float64(5), int64(5), object(16)
memory usage: 41.8+ KB

```

```
1 df = df.rename({"normalized-losses": "normalizedlosses"}, axis='columns')
```

```
1 df.info()
```

```

<class 'pandas.core.frame.DataFrame'>
RangeIndex: 205 entries, 0 to 204
Data columns (total 26 columns):
 #   Column      Non-Null Count  Dtype  
--- 
 0   symboling    205 non-null   int64  
 1   normalizedlosses  205 non-null  object  
 2   make         205 non-null   object  
 3   fuel-type    205 non-null   object  
 4   aspiration   205 non-null   object  
 5   num-of-doors 205 non-null   object  
 6   body-style   205 non-null   object  
 7   drive-wheels 205 non-null   object  
 8   engine-location 205 non-null  object  
 9   wheel-base   205 non-null   float64 
 10  length       205 non-null   float64 
 11  width        205 non-null   float64 
 12  height       205 non-null   float64 
 13  curb-weight  205 non-null   int64  
 14  engine-type  205 non-null   object  
 15  num-of-cylinders 205 non-null  object  
 16  engine-size  205 non-null   int64  
 17  fuel-system  205 non-null   object  
 18  bore         205 non-null   object  
 19  stroke       205 non-null   object  
 20  compression-ratio 205 non-null  float64 
 21  horsepower   205 non-null   object  
 22  peak-rpm     205 non-null   object  
 23  city-mpg     205 non-null   int64  
 24  highway-mpg  205 non-null   int64  
 25  price        205 non-null   object  
dtypes: float64(5), int64(5), object(16)
memory usage: 41.8+ KB

```

```
1 df.normalizedlosses.unique()
```

```

array(['?', '164', '158', '192', '188', '121', '98', '81', '118', '148',
       '110', '145', '137', '101', '78', '106', '85', '107', '104', '113',
       '150', '129', '115', '93', '142', '161', '153', '125', '128',
       '122', '103', '168', '108', '194', '231', '119', '154', '74',
       '186', '83', '102', '89', '87', '77', '91', '134', '65', '197',
       '90', '94', '256', '95'], dtype=object)

```

```
1 df = df[df.normalizedlosses != "?"]
```

```
2 df.normalizedlosses = df.normalizedlosses.astype("float")
```

```
1 df.horsepower.unique()
```

```

array(['102', '115', '110', '140', '101', '121', '48', '70', '68', '88',
       '145', '58', '76', '60', '86', '100', '176', '135', '84', '120',
       '123', '155', '116', '69', '55', '97', '152', '160', '200', '95',
       '142', '143', '73', '82', '94', '111', '62', '56', '112', '92',
       '161', '156', '52', '85', '90', '114', '162', '134', '106'],
      dtype=object)

```

```
1 df = df[df.horsepower != "?"]
```

```
2 df.horsepower = df.horsepower.astype("int")
```

```
1 df.stroke.unique()
```

```

array(['3.4', '2.8', '3.19', '3.03', '3.11', '3.23', '3.39', '3.46',
       '3.9', '3.41', '3.07', '3.58', '4.17', '3.15', '?', '3.16', '3.64'],
      dtype=object)

```

```
'3.1', '3.29', '3.47', '3.27', '3.52', '2.19', '3.21', '2.07',
'2.36', '2.64', '3.35', '3.08', '3.5', '3.54', '2.87'],
dtype=object)
```

```
1 df = df[df.stroke != "?"]
2 df.stroke = df.stroke.astype("float")

1 df = df.rename({"peak-rpm": "peakrpm"}, axis='columns')

1 df.info()
```

```
<class 'pandas.core.frame.DataFrame'>
Int64Index: 160 entries, 3 to 204
Data columns (total 26 columns):
 #   Column           Non-Null Count  Dtype  
--- 
 0   symboling        160 non-null    int64  
 1   normalizedlosses 160 non-null    float64 
 2   make             160 non-null    object  
 3   fuel-type        160 non-null    object  
 4   aspiration       160 non-null    object  
 5   num-of-doors     160 non-null    object  
 6   body-style       160 non-null    object  
 7   drive-wheels     160 non-null    object  
 8   engine-location   160 non-null    object  
 9   wheel-base       160 non-null    float64 
 10  length           160 non-null    float64 
 11  width            160 non-null    float64 
 12  height           160 non-null    float64 
 13  curb-weight      160 non-null    int64  
 14  engine-type      160 non-null    object  
 15  num-of-cylinders 160 non-null    object  
 16  engine-size      160 non-null    int64  
 17  fuel-system      160 non-null    object  
 18  bore              160 non-null    object  
 19  stroke            160 non-null    float64 
 20  compression-ratio 160 non-null    float64 
 21  horsepower        160 non-null    int32  
 22  peakrpm           160 non-null    object  
 23  city-mpg          160 non-null    int64  
 24  highway-mpg       160 non-null    int64  
 25  price              160 non-null    object  
dtypes: float64(7), int32(1), int64(5), object(13)
memory usage: 33.1+ KB
```

```
1 df.peakrpm.unique()

array(['5500', '5800', '4250', '5100', '5400', '5000', '4800', '6000',
'4750', '4350', '5200', '4150', '5600', '5250', '4900', '4400',
'4500', '6600', '4200', '5300'], dtype=object)
```

```
1 df = df[df.peakrpm != "?"]
2 df.peakrpm = df.peakrpm.astype("int")
```

```
1 df.price.unique()

array(['13950', '17450', '17710', '23875', '16430', '16925', '20970',
'21105', '5151', '6295', '6575', '5572', '6377', '7957', '6229',
'6692', '7609', '8558', '8921', '12964', '6479', '6855', '5399',
'6529', '7129', '7295', '7895', '9095', '8845', '10295', '12945',
'10345', '32250', '5195', '6095', '6795', '6695', '7395', '8495',
'10595', '10245', '11245', '18280', '25552', '28248', '28176',
'31600', '35056', '5389', '6189', '6669', '7689', '9959', '8499',
'6989', '8189', '9279', '5499', '7099', '6649', '6849', '7349',
'7299', '7799', '7499', '7999', '8249', '8949', '9549', '13499',
'14399', '17199', '19699', '18399', '11900', '13200', '15580',
'16900', '16630', '17950', '18150', '22018', '11850', '12170',
'15040', '15510', '18620', '5118', '7053', '7603', '7126', '7775',
'9960', '9233', '11259', '7463', '10198', '8013', '11694', '5348',
'6338', '6488', '6918', '7898', '8778', '6938', '7198', '7788',
'7738', '8358', '9258', '8058', '8238', '9298', '9538', '8449',
'9639', '9989', '11199', '11549', '17669', '8948', '10698', '9988',
'10898', '11248', '16558', '15998', '15690', '7975', '7995',
'8195', '9495', '9995', '9980', '12940', '13415', '15985', '16515',
'18420', '18950', '16845', '19045', '21485', '22470', '22625'],
dtype=object)
```

```
1 df = df[df.price != "?"]
2 df.price = df.price.astype("int")
```

```
1 df.info()
```

```
<class 'pandas.core.frame.DataFrame'>
Int64Index: 160 entries, 3 to 204
Data columns (total 26 columns):
 #   Column           Non-Null Count  Dtype  
--- 
 0   symboling        160 non-null    int64  
 1   normalizedlosses 160 non-null    float64 
 2   make             160 non-null    object  
 3   fuel-type        160 non-null    object  
 4   aspiration       160 non-null    object  
 5   num-of-doors     160 non-null    object  
 6   body-style       160 non-null    object  
 7   drive-wheels     160 non-null    object  
 8   engine-location   160 non-null    object  
 9   wheel-base       160 non-null    float64 
 10  length           160 non-null    float64 
 11  width            160 non-null    float64 
 12  height           160 non-null    float64 
 13  curb-weight      160 non-null    int64  
 14  engine-type      160 non-null    object  
 15  num-of-cylinders 160 non-null    object  
 16  engine-size      160 non-null    int64  
 17  fuel-system      160 non-null    object  
 18  bore              160 non-null    object  
 19  stroke            160 non-null    float64 
 20  compression-ratio 160 non-null    float64 
 21  horsepower        160 non-null    int32  
 22  peakrpm           160 non-null    int32  
 23  city-mpg          160 non-null    int64  
 24  highway-mpg       160 non-null    int64  
 25  price             160 non-null    int32  
dtypes: float64(7), int32(3), int64(5), object(11)
memory usage: 31.9+ KB
```

```
1 df.describe()
```

```
2
```

	symboling	normalizedlosses	wheel-base	length	width	height	?
count	160.000000	160.000000	160.000000	160.000000	160.000000	160.000000	160.000000
mean	0.737500	121.300000	98.235625	172.319375	65.596250	53.878750	2459.2
std	1.189511	35.602417	5.163763	11.548860	1.946999	2.276608	480.1
min	-2.000000	65.000000	86.600000	141.100000	60.300000	49.400000	1488.1
25%	0.000000	94.000000	94.500000	165.525000	64.000000	52.000000	2073.1
50%	1.000000	114.000000	96.900000	172.200000	65.400000	54.100000	2338.1
75%	2.000000	148.000000	100.600000	177.800000	66.500000	55.500000	2808.1

```
1 df.isnull().sum()
```

```
symboling      0
normalizedlosses 0
make          0
fuel-type      0
aspiration     0
num-of-doors   0
body-style     0
drive-wheels   0
engine-location 0
wheel-base     0
length         0
width          0
height         0
curb-weight    0
engine-type    0
num-of-cylinders 0
engine-size    0
```

```

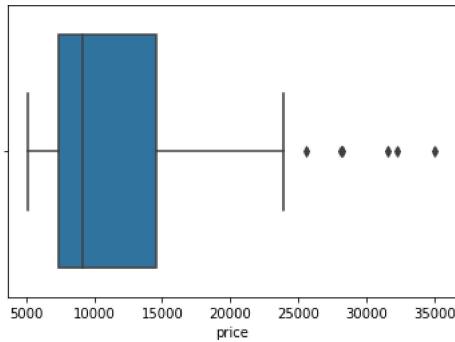
fuel-system      0
bore            0
stroke          0
compression-ratio 0
horsepower      0
peakrpm         0
city-mpg        0
highway-mpg     0
price           0
dtype: int64

```

```

1 import seaborn as sns
2 sns.boxplot(df["price"])
3
E:\Anaconda\lib\site-packages\seaborn\_decorators.py:36: FutureWarning: Pass the followi
  warnings.warn(
<AxesSubplot:xlabel='price'>

```



```

1 def outliers_iqr(a):
2     z=[]
3     quartile_1,quartile_3=np.percentile(a,[25,75])
4     iqr=quartile_3-quartile_1
5     lower_bound=quartile_1-(iqr*1.5)
6     upper_bound=quartile_3+(iqr*1.5)
7     x=np.where((a>upper_bound)|(a<lower_bound))
8     for i in x:
9         z.append(x)
10    return(z)

```

```

1 outliers_iqr(df["price"])
[(array([33, 45, 46, 47, 48, 49], dtype=int64),)]

```

```

1 df.drop([32, 47],axis=0,inplace=True)

```

```

1 df.head()

```

	symboling	normalizedlosses	make	fuel-type	aspiration	num-of-doors	body-style	drive-wheels	engine-location
3	2	164.0	audi	gas	std	four	sedan	fwd	front
4	2	164.0	audi	gas	std	four	sedan	4wd	front
6	1	158.0	audi	gas	std	four	sedan	fwd	front
8	1	158.0	audi	gas	turbo	four	sedan	fwd	front
10	2	192.0	bmw	gas	std	two	sedan	rwd	front

```

1 df.dtypes
2

```

```

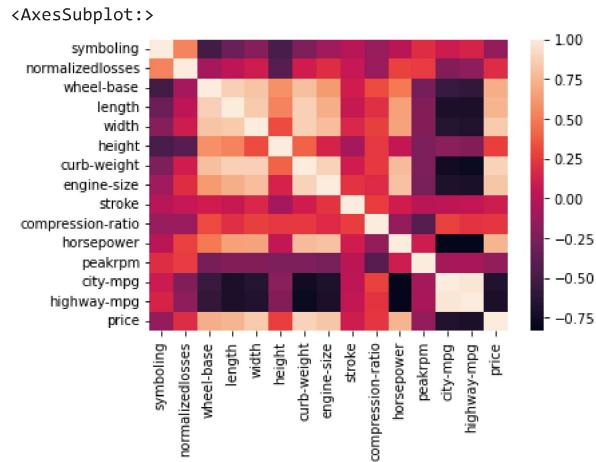
symboling      int64
normalizedlosses   float64
make          object
fuel-type     object
aspiration    object
num-of-doors  object
body-style    object
drive-wheels  object
engine-location  object
wheel-base    float64
length        float64
width         float64
height        float64
curb-weight   int64
engine-type   object
num-of-cylinders  object
engine-size   int64
fuel-system   object
bore          object
stroke        float64
compression-ratio  float64
horsepower    int32
peakrpm       int32
city-mpg      int64
highway-mpg   int64
price         int32
dtype: object

```

```

1 import seaborn as sns
2 sns.heatmap(df.corr())

```



```
1 df.corr()
```

```

symboling  normalizedlosses  wheel-
base      length      width      height
1 df=df.drop(["wheel-base"],axis=1)
2 df.head()
3 df.shape
(158, 25)

1 df=df.drop(["make"],axis=1)
2 df.head()
3 df.shape
(158, 24)

1 df=df.drop(["fuel-system","aspiration","num-of-doors","body-style","drive-wheels","engine-location","fuel-s
2 df.head()
3 df.shape
(158, 18)

1 df=df.drop(["fuel-type"],axis=1)
2 df.head()
3 df.shape
(158, 17)

1 df=df.drop(["engine-type"],axis=1)
2 df.head()
3 df.shape
(158, 16)

1 df=df.drop(["num-of-cylinders"],axis=1)
2 df.head()
3 df.shape
(158, 15)

1 df.head()

      symboling  normalizedlosses  length  width  height  curb-
      weight  engine-
      size   bore  stroke  c
3         2          164.0     176.6    66.2     54.3    2337      109    3.19    3.4
4         2          164.0     176.6    66.4     54.3    2824      136    3.19    3.4
6         1          158.0     192.7    71.4     55.7    2844      136    3.19    3.4
8         1          158.0     192.7    71.4     55.9    3086      131    3.13    3.4

```

1 from sklearn.decomposition import PCA
2 pca = PCA(n_components=1)
3 pca
PCA(n_components=1)

1 principalComponents = pca.fit_transform(df)
2

1 principalComponents.shape
(158, 1)

```
1 print(principalComponents[:])  
2
```

```
[[ 2595.28031554]  
[ 6121.21712957]  
[ 6381.95115049]  
[12547.68375819]  
[ 5068.73388433]  
[ 5562.37248451]  
[ 9639.9412759 ]  
[ 9778.56675541]  
[-6236.47886471]  
[-5071.32507407]  
[-4789.56777207]  
[-5793.49308563]  
[-4990.70048048]  
[-3396.55508139]  
[-5131.64464088]  
[-4668.31937731]  
[-3753.84152692]  
[-2792.5940468 ]  
[-2398.95706433]  
[ 1653.39275246]  
[-4891.62328263]  
[-4524.72507763]  
[-4841.06637537]  
[-4241.55393477]  
[-4072.05531666]  
[-4071.08532144]  
[-3454.51933067]  
[-2253.96840302]  
[-2502.20400175]  
[-1051.25142549]  
[ 1598.28427912]  
[-1003.08708259]  
[-6161.82398094]  
[-5263.57897355]  
[-4565.1409342 ]  
[-4661.9368335 ]  
[-3963.49879291]  
[-2482.99527051]  
[-2830.23243486]  
[ -737.80853313]  
[-1085.04569748]  
[ -86.70610008]  
[ 6944.3287484 ]  
[14266.51442304]  
[16972.17667301]  
[16881.83860748]  
[20316.42824154]  
[23751.82641139]  
[-5972.87774324]  
[-5173.18199344]  
[-4690.14286989]  
[-3662.53109573]  
[-1382.30422828]  
[-2834.79307445]  
[-4337.98529516]  
[-3138.37906294]  
[-2058.0699862 ]  
[-2058.05554879]
```

```
1 pca.explained_variance_ratio_
```

```
array([0.99192393])
```

```
1 plt.plot(range(2), pca.explained_variance_ratio_,label="Variance-ratio")  
2 plt.plot(range(2), np.cumsum(pca.explained_variance_ratio_),label="cumulative variance-ratio")  
3 plt.title("Component-wise and Cumulative Explained Variance")  
4 plt.xlabel("No.of components")  
5 plt.ylabel("Data_explained")
```

```

-----
ValueError
Input In [189], in <cell line: 1>()
----> 1 plt.plot(range(2), pca.explained_variance_ratio_, label="Variance-ratio")
      2 plt.plot(range(2), np.cumsum(pca.explained_variance_ratio_), label="cumulative
      variance-ratio")
      3 plt.title("Component-wise and Cumulative Explained Variance")

File E:\Anaconda\lib\site-packages\matplotlib\pyplot.py:2757, in plot(scaledx, scaledy,
data, *args, **kwargs)
    2755 @_copy_docstring_and_deprecators(Axes.plot)
    2756 def plot(*args, scaledx=True, scaledy=True, data=None, **kwargs):
--> 2757     return gca().plot(
    2758         *args, scaledx=scaledx, scaledy=scaledy,
    2759         **({ "data": data} if data is not None else {}), **kwargs)

File E:\Anaconda\lib\site-packages\matplotlib\axes\_axes.py:1632, in Axes.plot(self,
scaledx, scaledy, data, *args, **kwargs)
    1390 """
    1391 Plot y versus x as lines and/or markers.
    1392
    (...)

1629 (``'green'``) or hex strings (``'#008000'``).
1630 """
1631 kwargs = cbook.normalize_kwargs(kwargs, mlines.Line2D)
--> 1632 lines = [*self._get_lines(*args, data=data, **kwargs)]
1633 for line in lines:
1634     self.add_line(line)

File E:\Anaconda\lib\site-packages\matplotlib\axes\_base.py:312, in
_process_plot_var_args.__call__(self, data, *args, **kwargs)
    310     this += args[0],
    311     args = args[1:]
--> 312 yield from self._plot_args(this, kwargs)

File E:\Anaconda\lib\site-packages\matplotlib\axes\_base.py:498, in
_process_plot_var_args._plot_args(self, tup, kwargs, return_kwargs)
    495     self.axes.yaxis.update_units(y)
    497 if x.shape[0] != y.shape[0]:
--> 498     raise ValueError(f"x and y must have same first dimension, but "
    499                     f"have shapes {x.shape} and {y.shape}")
    500 if x.ndim > 2 or y.ndim > 2:
    501     raise ValueError(f"x and y can be no greater than 2D, but have "
    502                     f"shapes {x.shape} and {y.shape}")

ValueError: x and y must have same first dimension, but have shapes (2,) and (1,)

SEARCH STACK OVERFLOW

```



```

1 #Linear Regression
2 import numpy as np
3 import pandas as pd
4 import matplotlib.pyplot as plt
5 from sklearn.model_selection import train_test_split
6

1 from sklearn.metrics import mean_squared_error, r2_score

1 training_features = [ 'length','width','curb-weight','horsepower']
2 target = 'price'

1 X_train, X_test, Y_train, Y_test = train_test_split(df[training_features],
2                                         df[target],
3                                         test_size=0.1,random_state=222)

1 from sklearn import linear_model
2 model = linear_model.LinearRegression()
3 model.fit(X_train,Y_train)

```

```

1 model.coef_
2

1 model.intercept_

1 predicted=model.predict(X_test)
2 print("Mean squared error: %.2f"
3           % mean_squared_error(Y_test, predicted))

1 print('R Square score: %.2f' % r2_score(Y_test, predicted))

1 from sklearn.model_selection import cross_val_score

1 scores = cross_val_score(model, X_train, Y_train, cv=10)
2

1 scores

1 print("Accuracy: %0.2f (+/- %0.2f)" % (scores.mean(), 3 * scores.std() ))
Accuracy: 0.73 (+/- 0.30)

1 #Lasso Regression
2
3
4 model = linear_model.Lasso(alpha=5)
5 model.fit(X_train, Y_train)
6
7 from sklearn.metrics import mean_squared_error, r2_score
8 training_features = [ 'length', 'width', 'curb-weight', 'horsepower']
9 target = 'price'
10 X_train, X_test, Y_train, Y_test = train_test_split(df[training_features],
11                                                 df[target],
12                                                 test_size=0.1, random_state=222)
13
14

1 predicted=model.predict(X_test)
2 print("Mean squared error: %.2f"
3           % mean_squared_error(Y_test, predicted))
4 print('R Square score: %.2f' % r2_score(Y_test, predicted))

Mean squared error: 5874511.49
R Square score: 0.80

1 from sklearn.model_selection import cross_val_score
2
3 scores = cross_val_score(model, X_train, Y_train, cv=10)
4
5 scores
6

array([0.85283967, 0.5496465 , 0.09993309, 0.82811057, 0.78175393,
       0.83087709, 0.75097281, 0.68421817, 0.71114213, 0.81003868])

1 print("Accuracy: %0.2f (+/- %0.2f)" % (scores.mean(), 3 * scores.std() ))
Accuracy: 0.69 (+/- 0.64)

1 #RidgeRegression
2

```

```

3
4 import numpy as np
5 import pandas as pd
6 import matplotlib.pyplot as plt
7 from sklearn.model_selection import train_test_split
8
9 from sklearn.metrics import mean_squared_error, r2_score
10 training_features = [ 'length','width','curb-weight','horsepower']
11 target = 'price'
12 X_train, X_test, Y_train, Y_test = train_test_split(df[training_features],
13                                         df[target],
14                                         test_size=0.1,random_state=222)
15
16

1 model =linear_model.Ridge(alpha=10)
2 model.fit(X_train,Y_train)

Ridge(alpha=10)

1 predicted=model.predict(X_test)
2 print("Mean squared error: %.2f"
3         % mean_squared_error(Y_test, predicted))
4
5 print('R Square score: %.2f' % r2_score(Y_test, predicted))

Mean squared error: 6109120.80
R Square score: 0.79

1
2 from sklearn.model_selection import cross_val_score
3
4 scores = cross_val_score(model, X_train, Y_train,cv=10)
5
6 scores
7

array([0.85029325, 0.52637566, 0.13488245, 0.82427246, 0.78136647,
       0.83054025, 0.76262181, 0.69171734, 0.71140007, 0.80903254])

1 print("Accuracy: %0.2f (+/- %0.2f)" % (scores.mean(), 3* scores.std() ))
Accuracy: 0.69 (+/- 0.62)

1 #KNN Regressor
2
3
4 from sklearn.neighbors import KNeighborsRegressor
5 model = KNeighborsRegressor(n_neighbors=4,
6                             p=2,weights='distance')
7 model.fit(X_train,Y_train)

KNeighborsRegressor(n_neighbors=4, weights='distance')

1 predicted=model.predict(X_test)
2 print("Mean squared error: %.2f"
3         % mean_squared_error(Y_test, predicted))
4
5 print('R Square score: %.2f' % r2_score(Y_test, predicted))

Mean squared error: 12388238.85
R Square score: 0.58

```



```
E:\Anaconda\lib\site-packages\sklearn\neural_network\_multilayer_perceptron.py:692: ConvergenceWarning: Stochastic Optimizer: Maximum iteration limit reached without convergence. Consider increasing the `max_iter` parameter.
```

```
E:\Anaconda\lib\site-packages\sklearn\neural_network\_multilayer_perceptron.py:692: ConvergenceWarning: Stochastic Optimizer: Maximum iteration limit reached without convergence. Consider increasing the `max_iter` parameter.
```

```
array([-4.72904487, -6.76958747, -13.35388065, -2.4503772 ,  
       -3.35492748, -5.91468527, -4.16524559, -5.31747154,  
       -6.16531129, -3.29268685])
```

```
1  
2 print("Accuracy: %0.2f (+/- %0.2f)" % (scores.mean(), 3 * scores.std() ))  
3  
Accuracy: -5.55 (+/- 8.75)
```

```
1  
2  
3 #Regression trees  
4  
5 from sklearn.model_selection import train_test_split, GridSearchCV  
6  
7 from sklearn import datasets, linear_model  
8 from sklearn.metrics import mean_squared_error, r2_score  
9 from sklearn.model_selection import train_test_split  
10 from sklearn.tree import DecisionTreeRegressor  
11 model = DecisionTreeRegressor(criterion='mae', max_depth=5, min_samples_leaf=4, min_samples_split=4)  
12 model.fit(X_train, Y_train)
```

```
E:\Anaconda\lib\site-packages\sklearn\tree\_classes.py:366: FutureWarning: Criterion 'mae' was deprecated in v1.0 and will be removed in v1.1. Use 'mse' instead.
```

```
DecisionTreeRegressor(criterion='mae', max_depth=5, min_samples_leaf=4,  
                      min_samples_split=4)
```

```
1 predicted=model.predict(X_test)  
2 print("Mean squared error: %.2f"  
3      % mean_squared_error(Y_test, predicted))  
4  
5 print('R Square score: %.2f' % r2_score(Y_test, predicted))
```

```
Mean squared error: 9519435.94  
R Square score: 0.68
```

```
1 from sklearn.tree import plot_tree  
2 plot_tree(model,filled=True)
```

```
[Text(0.6176470588235294, 0.9166666666666666, 'X[2] <= 2664.0\nmae = 3818.239\nsamples = 142\nvalue = 8948.5'),
Text(0.3676470588235294, 0.75, 'X[2] <= 2291.5\nmae = 1502.549\nsamples = 102\nvalue = 7896.5'),
Text(0.19117647058823528, 0.5833333333333333, 'X[2] <= 1998.5\nmae = 835.453\nsamples = 64\nvalue = 7163.5'),
Text(0.11764705882352941, 0.4166666666666667, 'X[2] <= 1902.5\nmae = 558.679\nsamples = 28\nvalue = 6504.0'),
Text(0.058823529411764705, 0.25, 'X[2] <= 1875.0\nmae = 509.889\nsamples = 9\nvalue = 6095.0'),
Text(0.029411764705882353, 0.0833333333333333, 'mae = 472.0\nsamples = 4\nvalue = 6387.0'),
Text(0.08823529411764706, 0.0833333333333333, 'mae = 355.6\nsamples = 5\nvalue = 5572.0'),
Text(0.17647058823529413, 0.25, 'X[0] <= 158.75\nmae = 487.0\nsamples = 19\nvalue = 6692.0'),
Text(0.14705882352941177, 0.0833333333333333, 'mae = 546.1\nsamples = 10\nvalue = 6229.0'),
Text(0.20588235294117646, 0.0833333333333333, 'mae = 270.111\nsamples = 9\nvalue = 6795.0'),
Text(0.2647058823529412, 0.4166666666666667, 'X[1] <= 63.7\nmae = 631.056\nsamples = 36\nvalue = 7775.0'),
Text(0.23529411764705882, 0.25, 'mae = 503.0\nsamples = 5\nvalue = 6488.0'),
Text(0.29411764705882354, 0.25, 'X[3] <= 83.5\nmae = 506.71\nsamples = 31\nvalue = 7799.0'),
Text(0.2647058823529412, 0.0833333333333333, 'mae = 388.667\nsamples = 21\nvalue = 7738.0'),
Text(0.3235294117647059, 0.0833333333333333, 'mae = 571.5\nsamples = 10\nvalue = 8345.0'),
Text(0.5441176470588235, 0.5833333333333334, 'X[2] <= 2422.5\nmae =
```

```
1 from sklearn.model_selection import cross_val_score
2
3 scores = cross_val_score(model, X_train, Y_train, cv=10)
4
5 scores
```

```
E:\Anaconda\lib\site-packages\sklearn\tree\_classes.py:366: FutureWarning: Criterion 'mae' was deprecated in v1.0 and will be removed in
warnings.warn(
E:\Anaconda\lib\site-packages\sklearn\tree\_classes.py:366: FutureWarning: Criterion 'mae' was deprecated in v1.0 and will be removed in
warnings.warn(
E:\Anaconda\lib\site-packages\sklearn\tree\_classes.py:366: FutureWarning: Criterion 'mae' was deprecated in v1.0 and will be removed in
warnings.warn(
E:\Anaconda\lib\site-packages\sklearn\tree\_classes.py:366: FutureWarning: Criterion 'mae' was deprecated in v1.0 and will be removed in
warnings.warn(
E:\Anaconda\lib\site-packages\sklearn\tree\_classes.py:366: FutureWarning: Criterion 'mae' was deprecated in v1.0 and will be removed in
warnings.warn(
E:\Anaconda\lib\site-packages\sklearn\tree\_classes.py:366: FutureWarning: Criterion 'mae' was deprecated in v1.0 and will be removed in
warnings.warn(
E:\Anaconda\lib\site-packages\sklearn\tree\_classes.py:366: FutureWarning: Criterion 'mae' was deprecated in v1.0 and will be removed in
warnings.warn(
E:\Anaconda\lib\site-packages\sklearn\tree\_classes.py:366: FutureWarning: Criterion 'mae' was deprecated in v1.0 and will be removed in
warnings.warn(
E:\Anaconda\lib\site-packages\sklearn\tree\_classes.py:366: FutureWarning: Criterion 'mae' was deprecated in v1.0 and will be removed in
warnings.warn(
E:\Anaconda\lib\site-packages\sklearn\tree\_classes.py:366: FutureWarning: Criterion 'mae' was deprecated in v1.0 and will be removed in
warnings.warn(
array([0.81474477, 0.62690434, 0.44097421, 0.82795639, 0.89808225,
       0.74624446, 0.88347892, 0.71621996, 0.77092197, 0.94997441])
```

```
1 print("Accuracy: %0.2f (+/- %0.2f)" % (scores.mean(), 3 * scores.std() ))
```

```
Accuracy: 0.77 (+/- 0.42)
```

```
1 #SVR
2
3 from sklearn.svm import SVR
4 model = SVR(kernel='linear')
5 model.fit(X_train, Y_train)

SVR(kernel='linear')

1 predicted=model.predict(X_test)
2 print("Mean squared error: %.2f"
3      % mean_squared_error(Y_test, predicted))
```

```
4
5 print('R Square score: %.2f' % r2_score(Y_test, predicted))

Mean squared error: 9794854.40
R Square score: 0.67

1
2 from sklearn.model_selection import cross_val_score
3
4 scores = cross_val_score(model, X_train, Y_train, cv=10)
5
6 scores

array([0.81111938, 0.6505549 , 0.48148009, 0.73129806, 0.73213532,
       0.80326901, 0.81236196, 0.68020696, 0.81925003, 0.76819261])

1 print("Accuracy: %0.2f (+/- %0.2f)" % (scores.mean(), 3 * scores.std() ))

Accuracy: 0.73 (+/- 0.30)

1
1
1
1
1
```

✓ 1s completed at 4:05 PM

